Proofs for Smoothness of Parametric Regression Models

November 4, 2016

Intro

In this document, we consider parametric regression models $g(\cdot|\boldsymbol{\theta})$ where $\boldsymbol{\theta} \in \mathbb{R}^p$. Throughout, we will suppose $\boldsymbol{\theta}^*$ is the model such that

$$\boldsymbol{\theta}^* = \arg\min_{\theta \in \Theta} E_{x,y}\left[(y - g(x|\boldsymbol{\theta}))^2\right]$$

Technically, all the proofs require is that $\boldsymbol{\theta}^* \in \Theta$ is fixed. In the convergence rate proofs, we will need $\boldsymbol{\theta}^*$ to satisfy $E[y|x] = g(x|\boldsymbol{\theta}^*)$.
We are interested in establishing inequalities of the form

$$\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} - \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(2)}}\|_2 \leq C\|\boldsymbol{\lambda}^{(2)} - \boldsymbol{\lambda}^{(1)}\|_2$$

If the functions are Lipschitz in their parameterization, we will also be able to bound the distance between the actual functions. That is, if there are constants $L > 0$ and $r \in \mathbb{R}$, such that for all $\boldsymbol{\theta_1}, \boldsymbol{\theta_2}$

$$\|g(\cdot|\boldsymbol{\theta}_1) - g(\cdot|\boldsymbol{\theta}_2)\|_\infty \leq Lp^r\|\boldsymbol{\theta}_1 - \boldsymbol{\theta}_2\|_2$$

Then

$$\|g(\cdot|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) - g(\cdot|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(2)}})\|_\infty \leq Lp^rC\|\boldsymbol{\lambda}^{(2)} - \boldsymbol{\lambda}^{(1)}\|_2$$

Document Outline

First, we consider smooth training criteria and prove smoothness for two parametric regression examples:

1. Multiple penalties for a single model

$$\hat{\boldsymbol{\theta}}_\lambda = \arg\min_{\theta \in \mathbb{R}^p} \frac{1}{2}\|y - g(\cdot|\boldsymbol{\theta})\|_T^2 + \sum_{j=1}^{J} \lambda_j \left(P_j(\boldsymbol{\theta}) + \frac{w}{2}\|\boldsymbol{\theta}\|_2^2\right)$$

2. Additive model

$$\hat{\boldsymbol{\theta}}_\lambda = \arg\min_{\theta\in\mathbb{R}^p} \frac{1}{2}\|y - \sum_{j=1}^{J} g_j(\cdot|\boldsymbol{\theta}_j)\|_T^2 + \sum_{j=1}^{J} \lambda_j \left(P_j(\boldsymbol{\theta}_j) + \frac{w}{2}\|\boldsymbol{\theta}_j\|_2^2\right)$$

Then we will extend these results to non-smooth penalty functions.

Finally we will consider examples of parametric penalty functions. This includes a deep dive into the Sobolev penalty.

1 Multiple smooth penalties for a single model

The function class of interest are the minimizers of the penalized least squares criterion:

$$\mathcal{G}(T) = \left\{ \hat{\boldsymbol{\theta}}_\lambda = \arg\min_{\theta\in\mathbb{R}^p} \frac{1}{2}\|y - g(\cdot|\boldsymbol{\theta})\|_T^2 + \sum_{j=1}^{J} \lambda_j \left(P_j(\boldsymbol{\theta}) + \frac{w}{2}\|\boldsymbol{\theta}\|_2^2\right) : \boldsymbol{\lambda} \in \Lambda \right\}$$

where $\Lambda = [\lambda_{min}, \lambda_{max}]^J$ and $w > 0$ is a fixed constant.

Suppose that the penalties and the function $g(x|\boldsymbol{\theta})$ are twice-differentiable and convex wrt $\boldsymbol{\theta}$:

- Suppose that $\nabla_\theta^2 P_j(\boldsymbol{\theta})$ are PSD matrices for all $j = 1, ..., J$.
- Suppose that $\nabla_\theta^2 \|y - g(x|\boldsymbol{\theta})\|_T^2$ is a PSD matrix.

Suppose there is some $K > 0$ such that for all $j = 1, ..., J$ and any $\boldsymbol{\theta}, \boldsymbol{\beta}, m'$, we have

$$\left| \frac{\partial}{\partial m} P_j\left(\boldsymbol{\theta} + m\boldsymbol{\beta}\right)\Big|_{m=m'} \right| \leq K\|\boldsymbol{\beta}\|_2$$

Then for any $\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}} \in \Lambda$ we have

$$\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} - \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(2)}}\|_2 \leq \|\boldsymbol{\lambda^{(2)}} - \boldsymbol{\lambda^{(1)}}\|_2 \left(w\sqrt{J}\lambda_{min}\right)^{-1} \left(K + w\sqrt{\frac{2}{J\lambda_{min}w}\left(1 + \frac{J\lambda_{max}}{\lambda_{min}}\right) C_J}\right)$$

where

$$C_J = \frac{1}{2}\|y - g(\cdot|\boldsymbol{\theta}^*)\|_T^2 + \lambda_{max} \sum_{j=1}^{J} \left(P_j(\boldsymbol{\theta}^*) + \frac{w}{2}\|\boldsymbol{\theta}^*\|_2^2\right)$$

(Still thinking about how to make C_J more clear. Perhaps I can find a nicer bound for the Lipschitz constant)

Proof

Consider any $\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}} \in \Lambda$. Let $\boldsymbol{\beta} = \hat{\boldsymbol{\theta}}_{\lambda^{(2)}} - \hat{\boldsymbol{\theta}}_{\lambda^{(1)}}$.

Define

$$\hat{m}_{\boldsymbol{\beta}}(\boldsymbol{\lambda}) = \arg\min_{m \in \mathbb{R}} \frac{1}{2} \left\| y - g(\cdot | \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j \left(P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}\|_2^2 \right)$$

By definition, we know that $\hat{m}_{\beta}(\boldsymbol{\lambda^{(2)}}) = 1$ and $\hat{m}_{\beta}(\boldsymbol{\lambda^{(1)}}) = 0$.

1. We calculate $\nabla_{\lambda} \hat{m}_{\boldsymbol{\beta}}(\boldsymbol{\lambda})$ using the implicit differentiation trick.

By the KKT conditions, we have

$$\frac{\partial}{\partial m} \left(\frac{1}{2} \left\| y - g(\cdot | \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right) + \sum_{j=1}^{J} \lambda_j w \langle \boldsymbol{\beta}, \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta} \rangle \Bigg|_{m = \hat{m}_{\beta}(\boldsymbol{\lambda})} = 0$$

Now we implicitly differentiate with respect to λ_ℓ for $\ell = 1, 2, ..., J$

$$\frac{\partial}{\partial \lambda_\ell} \left\{ \left[\frac{\partial}{\partial m} \left(\frac{1}{2} \left\| y - g(\cdot | \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right) + \sum_{j=1}^{J} \lambda_j w \langle \boldsymbol{\beta}, \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta} \rangle \right] \Bigg|_{m = \hat{m}_{\beta}(\boldsymbol{\lambda})} \right\} = 0$$

By the product rule and chain rule, we have

$$\left\{ \left[\frac{\partial^2}{\partial m^2} \left(\frac{1}{2} \left\| y - g(\cdot | \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right) + \sum_{j=1}^{J} \lambda_j w \|\boldsymbol{\beta}\|_2^2 \right] \frac{\partial}{\partial \lambda_\ell} \hat{m}_{\beta}(\boldsymbol{\lambda}) + \frac{\partial}{\partial m} P_\ell(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) + w \langle \boldsymbol{\beta}, \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta} \rangle \right\} \Bigg|_{m = \hat{m}_{\beta}(\boldsymbol{\lambda})} = 0$$

Rearranging, for every $\ell = 1, ..., J$, we get

$$\frac{\partial}{\partial \lambda_\ell} \hat{m}_{\boldsymbol{\beta}}(\boldsymbol{\lambda}) = - \left[\frac{\partial^2}{\partial m^2} \left(\frac{1}{2} \left\| y - g(\cdot | \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right) + \sum_{j=1}^{J} \lambda_j w \|\boldsymbol{\beta}\|_2^2 \right]^{-1} \left[\frac{\partial}{\partial m} P_\ell(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) + w \langle \boldsymbol{\beta}, \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta} \rangle \right] \Bigg|_{m = \hat{m}_{\beta}(\boldsymbol{\lambda})}$$

In vector notation, we have

$$\nabla_{\lambda} \hat{m}_{\boldsymbol{\beta}}(\boldsymbol{\lambda}) = - \left[\frac{\partial^2}{\partial m^2} \left(\frac{1}{2} \left\| y - g(\cdot | \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) \right) + \sum_{j=1}^{J} \lambda_j w \|\boldsymbol{\beta}\|_2^2 \right]^{-1} \left[\nabla_m P(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta}) + w \langle \boldsymbol{\beta}, \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} + m\boldsymbol{\beta} \rangle \mathbf{1} \right] \Bigg|_{m = \hat{m}_{\beta}(\boldsymbol{\lambda})}$$

where $\nabla_m P(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + m\boldsymbol{\beta})$ is the J-dimensional vector

$$\nabla_m P(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + m\boldsymbol{\beta}) = \begin{bmatrix} \frac{\partial}{\partial m} P_1(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + m\boldsymbol{\beta}) \\ ... \\ \frac{\partial}{\partial m} P_J(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + m\boldsymbol{\beta}) \end{bmatrix}$$

2. Bound $\|\nabla_\lambda \hat{m}_{\boldsymbol{\beta}}(\boldsymbol{\lambda})\|$
Bounding the first multiplicand:
The first multiplicand is bounded by

$$\left| \frac{\partial^2}{\partial m^2} \left(\frac{1}{2} \left\| y - g(\cdot|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + m\boldsymbol{\beta}) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + m\boldsymbol{\beta}) \right) + \sum_{j=1}^{J} \lambda_j w \|\boldsymbol{\beta}\|_2^2 \right|^{-1} \leq \left(wJ\lambda_{min} \|\boldsymbol{\beta}\|_2^2 \right)^{-1}$$

since the mean squared error and the penalty functions are convex.
Bounding the second multiplicand:
The first summand in the second multiplicand is bounded by assumption

$$\left| \frac{\partial}{\partial m} P_\ell(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + m\boldsymbol{\beta}) \right| \leq K\|\boldsymbol{\beta}\|_2$$

The second summand in the second multiplicand is bounded by

$$\left| w\langle \boldsymbol{\beta}, \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + \hat{m}_\beta(\boldsymbol{\lambda})\boldsymbol{\beta} \rangle \right| \quad \leq \quad w\|\boldsymbol{\beta}\|_2 \|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + \hat{m}_\beta(\boldsymbol{\lambda})\boldsymbol{\beta}\|_2 \tag{1}$$

We need to bound $\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + \hat{m}_\beta(\boldsymbol{\lambda})\boldsymbol{\beta}\|_2$. By definition of $\hat{m}_\beta(\boldsymbol{\lambda})$,

$$\begin{aligned} \sum_{j=1}^{J} \lambda_j \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + \hat{m}_\beta(\boldsymbol{\lambda})\boldsymbol{\beta}\|_2^2 &\leq \frac{1}{2}\|y - g(\cdot|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}})\|_T^2 + \sum_{j=1}^{J} \lambda_j \left(P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}\|_2^2 \right) \\ &= \frac{1}{2}\|y - g(\cdot|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}})\|_T^2 + \sum_{j=1}^{J} \lambda_j^{(1)} \left(P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}\|_2^2 \right) + \sum_{j=1}^{J} \left(\lambda_j - \lambda_j^{(1)} \right) \left(P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}\|_2^2 \right) \end{aligned}$$

To bound the first part of the right hand side, use the definition of $\hat{\boldsymbol{\theta}}_{\lambda^{(1)}}$:

$$\begin{aligned}\frac{1}{2}\|y - g(\cdot|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}})\|_T^2 + \sum_{j=1}^{J} \lambda_j^{(1)} \left(P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}\|_2^2\right) &\leq \frac{1}{2}\|y - g(\cdot|\boldsymbol{\theta}^*)\|_T^2 + \sum_{j=1}^{J} \lambda_j^{(1)} \left(P_j(\boldsymbol{\theta}^*) + \frac{w}{2}\|\boldsymbol{\theta}^*\|_2^2\right) \\ &\leq \frac{1}{2}\|y - g(\cdot|\boldsymbol{\theta}^*)\|_T^2 + \lambda_{max} \sum_{j=1}^{J} \left(P_j(\boldsymbol{\theta}^*) + \frac{w}{2}\|\boldsymbol{\theta}^*\|_2^2\right) \\ &= C\end{aligned}$$

To bound the second part of the right hand side, note that

$$\begin{aligned}\sum_{j=1}^{J} \left(\lambda_j - \lambda_j^{(1)}\right) \left(P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}\|_2^2\right) &\leq \sum_{j=1}^{J} \left(\lambda_j - \lambda_j^{(1)}\right) \left[\max_{k=1:J} P_k(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)}}\|_2^2\right] \\ &\leq J\lambda_{max} \left[\max_{k=1:J} P_k(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)}}\|_2^2\right]\end{aligned}$$

Combining the above three inequalities, we get

$$\sum_{j=1}^{J} \lambda_j \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + \hat{m}_\beta(\boldsymbol{\lambda})\boldsymbol{\beta}\|_2^2 \leq C + J\lambda_{max} \left[\max_{k=1:J} P_k(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)}}\|_2^2\right] \tag{2}$$

To bound $\max_{k=1:J} P_k(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)}}\|_2^2$, we note that by the definition of $\hat{\boldsymbol{\theta}}_{\lambda^{(1)}}$, we have

$$\begin{aligned}\sum_{j=1}^{J} \lambda_j^{(1)} \left(P_j(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)}}\|_2^2\right) &\leq \frac{1}{2}\|y - g(\cdot|\boldsymbol{\theta}^*)\|_T^2 + \sum_{j=1}^{J} \lambda_j^{(1)} \left(P_j(\boldsymbol{\theta}^*) + \frac{w}{2}\|\boldsymbol{\theta}^*\|_2^2\right) \\ &\leq C\end{aligned}$$

Therefore

$$\max_{k=1:J} P_k(\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)}}\|_2^2 \leq \frac{C}{\lambda_{min}} \tag{3}$$

Plugging (3) into (2) above, we get

$$\sum_{j=1}^{J} \lambda_j \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}} + \hat{m}_\beta(\boldsymbol{\lambda})\boldsymbol{\beta}\|_2^2 \leq \left(1 + \frac{J\lambda_{max}}{\lambda_{min}}\right) C \tag{4}$$

We can combine (4) with the fact that

$$J\lambda_{min}\frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}+\hat{m}_{\beta}(\boldsymbol{\lambda})\boldsymbol{\beta}\|_2^2\leq\sum_{j=1}^{J}\lambda_j\frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}+\hat{m}_{\beta}(\boldsymbol{\lambda})\boldsymbol{\beta}\|_2^2$$

to get

$$\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}+\hat{m}_{\beta}(\boldsymbol{\lambda})\boldsymbol{\beta}\|_2\leq\sqrt{\frac{2}{J\lambda_{min}w}\left(1+\frac{J\lambda_{max}}{\lambda_{min}}\right)C}$$

Plug the inequality above into (1) to get

$$w\langle\boldsymbol{\beta},\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)}}+\hat{m}_{\beta}(\boldsymbol{\lambda})\boldsymbol{\beta}\rangle\leq w\|\boldsymbol{\beta}\|_2\sqrt{\frac{2}{J\lambda_{min}w}\left(1+\frac{J\lambda_{max}}{\lambda_{min}}\right)C}$$

Finally we have bounded the derivative of $\frac{\partial}{\partial\lambda_\ell}\hat{m}_{\beta}(\boldsymbol{\lambda})$. For every $\ell=1,...,J$, we have

$$\begin{aligned}\left|\frac{\partial}{\partial\lambda_\ell}\hat{m}_{\beta}(\boldsymbol{\lambda})\right| &\leq \left(wJ\lambda_{min}\|\boldsymbol{\beta}\|_2^2\right)^{-1}\left(K\|\boldsymbol{\beta}\|_2+w\|\boldsymbol{\beta}\|_2\sqrt{\frac{2}{J\lambda_{min}w}\left(1+\frac{J\lambda_{max}}{\lambda_{min}}\right)C}\right)\\ &= \left(wJ\lambda_{min}\|\boldsymbol{\beta}\|_2\right)^{-1}\left(K+w\sqrt{\frac{2}{J\lambda_{min}w}\left(1+\frac{J\lambda_{max}}{\lambda_{min}}\right)C}\right)\end{aligned}$$

We can sum up these bounds to bound the norm of the gradient $\nabla_\lambda\hat{m}_{\beta}(\boldsymbol{\lambda})$:

$$\begin{aligned}\|\nabla_\lambda\hat{m}_{\beta}(\boldsymbol{\lambda})\| &= \sqrt{\sum_{\ell=1}^{J}\left(\frac{\partial}{\partial\lambda_\ell}\hat{m}_{\beta}(\boldsymbol{\lambda})\right)^2}\\ &\leq \left(w\lambda_{min}\sqrt{J}\|\boldsymbol{\beta}\|_2\right)^{-1}\left(K+w\sqrt{\frac{2}{J\lambda_{min}w}\left(1+\frac{J\lambda_{max}}{\lambda_{min}}\right)C}\right)\end{aligned}$$

3. Apply Mean Value Theorem

Since the training criterion is smooth, then $\hat{m}_{\boldsymbol{\beta}}(\boldsymbol{\lambda})$ is continuous and differentiable over the line segment $\left\{\alpha\boldsymbol{\lambda}^{(1)}+(1-\alpha)\boldsymbol{\lambda}^{(2)}:\alpha\in[0,1]\right\}$.

Therefore by MVT, there is some $\alpha \in (0,1)$ such that

$$\begin{aligned}\left|\hat{m}_\beta(\boldsymbol{\lambda^{(2)}}) - \hat{m}_\beta(\boldsymbol{\lambda^{(1)}})\right| &= \left|\left\langle \boldsymbol{\lambda^{(2)}} - \boldsymbol{\lambda^{(1)}}, \nabla_\lambda \hat{m}_\beta(\boldsymbol{\lambda})\right\rangle\Big|_{\boldsymbol{\lambda}=\alpha\boldsymbol{\lambda^{(1)}}+(1-\alpha)\boldsymbol{\lambda^{(2)}}}\right| \\ &\leq \|\boldsymbol{\lambda^{(2)}} - \boldsymbol{\lambda^{(1)}}\|_2 \left\| \nabla_\lambda \hat{m}_\beta(\boldsymbol{\lambda})|_{\boldsymbol{\lambda}=\alpha\boldsymbol{\lambda^{(1)}}+(1-\alpha)\boldsymbol{\lambda^{(2)}}} \right\| \\ &\leq \|\boldsymbol{\lambda^{(2)}} - \boldsymbol{\lambda^{(1)}}\|_2 \left(w\sqrt{J}\lambda_{min}\|\boldsymbol{\beta}\|_2\right)^{-1} \left(K + w\sqrt{\frac{2}{J\lambda_{min}w}\left(1+\frac{J\lambda_{max}}{\lambda_{min}}\right)C}\right)\end{aligned}$$

Recall that $\hat{m}_\beta(\boldsymbol{\lambda^{(2)}}) - \hat{m}_\beta(\boldsymbol{\lambda^{(1)}}) = 1$. Rearranging, we get

$$\|\boldsymbol{\beta}\|_2 = \|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} - \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(2)}}}\|_2 \leq \|\boldsymbol{\lambda^{(2)}} - \boldsymbol{\lambda^{(1)}}\|_2 \left(w\sqrt{J}\lambda_{min}\right)^{-1} \left(K + w\sqrt{\frac{2}{J\lambda_{min}w}\left(1+\frac{J\lambda_{max}}{\lambda_{min}}\right)C}\right)$$

2 Additive Model

The function class of interest are the minimizers of the penalized least squares criterion:

$$\mathcal{G}(T) = \left\{ \hat{\boldsymbol{\theta}}_\lambda = \arg\min_{\theta\in\mathbb{R}^p} \frac{1}{2}\left\| y - \sum_{j=1}^J g_j(\cdot|\boldsymbol{\theta}_j)\right\|_T^2 + \sum_{j=1}^J \lambda_j\left(P_j(\boldsymbol{\theta}_j) + \frac{w}{2}\|\boldsymbol{\theta}_j\|_2^2\right) : \boldsymbol{\lambda}\in\Lambda \right\}$$

where $\Lambda = [\lambda_{min}, \lambda_{max}]^J$.
Suppose that the penalties and the mean squared error $\|y - \sum_{j=1}^J g_j(x|\boldsymbol{\theta}_j)\|_T^2$ are twice-differentiable and convex wrt $\boldsymbol{\theta}$

- $\nabla^2_{\boldsymbol{\theta}_j} P_j(\boldsymbol{\theta}_j)$ are PSD matrices for all $j = 1, ..., J$
- $\nabla^2_{\boldsymbol{\theta}}\|y - \sum_{j=1}^J g_j(x|\boldsymbol{\theta}_j)\|_T^2$ is a PSD matrix.

Suppose for each $j = 1, ..., J$, there is a constant $K_j \geq 0$ such that for all $\boldsymbol{\beta}, \boldsymbol{\theta}, m'$, we either have

$$\left|\frac{\partial}{\partial m} P_j(\boldsymbol{\theta} + m\boldsymbol{\beta})\right|_{m=m'} \leq K_j\|\boldsymbol{\beta}\|_2 \tag{5}$$

or

$$\left\| \frac{\partial}{\partial m} g_j(X_{T,j}|\boldsymbol{\theta} + m\boldsymbol{\beta})\Big|_{m=m'}\right\| = \sqrt{\sum_{i=1}^n \left(\frac{\partial}{\partial m} g_j(x_{T,j}|\boldsymbol{\theta} + m\boldsymbol{\beta})\Big|_{m=m'}\right)^2} \leq K_j\|\boldsymbol{\beta}\|_2 \tag{6}$$

(These conditions bound the spectrum of the penalty function or the function itself.)
Let

$$C = \frac{1}{2}\left\| y - \sum_{j=1}^{J} g_j(\cdot|\boldsymbol{\theta}_j^*)\right\|_T^2 + \lambda_{max}\sum_{j=1}^{J}\left(P_j(\boldsymbol{\theta}_j^*) + \frac{w}{2}\|\boldsymbol{\theta}_j^*\|_2^2\right)$$

For $j = 1, .., J$, let

$$d_j = \begin{cases} \left(K_j + w\sqrt{\left(1+\frac{J\lambda_{max}}{\lambda_{min}}\right)\frac{2C}{\lambda_{min}w}}\right) & \text{if assumption (5) holds for } P_j \\ \frac{1}{\lambda_{min}}K_j\sqrt{\left(1+\frac{J\lambda_{max}}{\lambda_{min}}\right)2C} & \text{if assumption (6) holds for } g_j \end{cases}$$

Then for any $\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}} \in \Lambda$ we have for all $j = 1, ..., J$

$$\|\boldsymbol{\theta}_{\lambda^{(1)},j} - \boldsymbol{\theta}_{\lambda^{(2)},j}\| \leq \left\|\boldsymbol{\lambda}^{(2)} - \boldsymbol{\lambda}^{(1)}\right\| \lambda_{min}^{-1} w^{-1}\left(\max_{j=1,..,J} d_j\right)$$

Proof

Consider any $\boldsymbol{\lambda}^{(1)}, \boldsymbol{\lambda}^{(2)} \in \Lambda$. Let $\boldsymbol{\beta}_j = \hat{\boldsymbol{\theta}}_{\lambda^{(2)},j} - \hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}$ for all $j = 1, ..., J$.
Define

$$\hat{\boldsymbol{m}}(\boldsymbol{\lambda}) = \arg\min_{\boldsymbol{m}} \frac{1}{2}\left\| y - \sum_{j=1}^{J} g_j(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j)\right\|_T^2 + \sum_{j=1}^{J}\lambda_j\left(P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j) + \frac{w}{2}\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j\|_2^2\right)$$

By definition, we know that $\hat{\boldsymbol{m}}(\boldsymbol{\lambda}^{(2)}) = \mathbf{1}$ and $\hat{\boldsymbol{m}}(\boldsymbol{\lambda^{(1)}}) = \mathbf{0}$.
1. We calculate $\nabla_\lambda \hat{\boldsymbol{m}}_{\boldsymbol{k}}(\boldsymbol{\lambda})$ using the implicit differentiation trick.
By the KKT conditions, we have for all $j = 1 : J$

$$\frac{\partial}{\partial m_j}\left(\frac{1}{2}\left\| y - \sum_{j=1}^{J} g_j(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j)\right\|_T^2 + \lambda_j P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j)\right) + \lambda_j w\langle\boldsymbol{\beta}_j, \hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j\rangle\Bigg|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})} = 0 \tag{7}$$

Now we implicitly differentiate with respect to λ_ℓ for $\ell = 1, 2, ..., J$

$$\frac{\partial}{\partial\lambda_\ell}\left\{\left[\frac{\partial}{\partial m_j}\left(\frac{1}{2}\left\| y - \sum_{j=1}^{J} g_j(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j)\right\|_T^2 + \lambda_j P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j)\right) + \lambda_j w\langle\boldsymbol{\beta}_j, \hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j\rangle\right]\Bigg|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})}\right\} = 0$$

By the product rule and chain rule, we have

$$\left\{\sum_{k=1}^{J}\left[\frac{\partial^2}{\partial m_k \partial m_j}\left(\frac{1}{2}\left\|y-\sum_{j=1}^{J} g_j(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}+m_j\boldsymbol{\beta}_j)\right\|_T^2+1[k=j]\lambda_j P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}+m_j\boldsymbol{\beta}_j)\right)+1[k=j]\lambda_j w\|\boldsymbol{\beta}_j\|_2^2\right]\frac{\partial}{\partial\lambda_\ell}\hat{m}_k(\boldsymbol{\lambda})\right\}\Bigg|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})}$$
$$+1[j=\ell]\left\{\frac{\partial}{\partial m_\ell}P_\ell(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},\ell}+m_\ell\boldsymbol{\beta}_\ell)+w\langle\boldsymbol{\beta}_\ell,\hat{\boldsymbol{\theta}}_{\lambda^{(1)},\ell}+m_\ell\boldsymbol{\beta}_\ell\rangle\right\}\Bigg|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})} = 0$$

Define the following matrices

$$S: S_{jk} = \frac{\partial^2}{\partial m_k \partial m_j}\frac{1}{2}\left\|y-\sum_{j=1}^{J} g_j(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}+m_j\boldsymbol{\beta}_j)\right\|_T^2\Bigg|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})}$$

$$D_1 = diag\left(\frac{\partial^2}{\partial m_j^2}\lambda_j P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}+m_j\boldsymbol{\beta}_j)\right)\Bigg|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})}$$

$$D_2 = diag\left(\lambda_j w\|\boldsymbol{\beta}_j\|_2^2\right)$$

$$D_3 = diag\left(\frac{\partial}{\partial m_\ell}P_\ell(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},\ell}+m_\ell\boldsymbol{\beta}_\ell)+w\langle\boldsymbol{\beta}_\ell,\hat{\boldsymbol{\theta}}_{\lambda^{(1)},\ell}+m_\ell\boldsymbol{\beta}_\ell\rangle\right)\Bigg|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})}$$

$$M = \begin{pmatrix} \nabla_\lambda\hat{m}_1(\lambda) & \nabla_\lambda\hat{m}_2(\lambda) & ... & \nabla_\lambda\hat{m}_J(\lambda) \end{pmatrix}$$

We can then combine all the equations into the following system of equations:

$$M = -D_3\left(S+D_1+D_2\right)^{-1}$$

S is a PSD matrix since the composition of a convex function with an affine function is convex.
D_1 is a PSD matrix since the penalty functions are convex.
2. We bound every diagonal element in D_3:
By Cauchy-Schwarz,

$$\left|w\langle\boldsymbol{\beta}_k,\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k}+\hat{m}_k(\boldsymbol{\lambda})\boldsymbol{\beta}_k\rangle\right| \le w\|\boldsymbol{\beta}_k\|\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k}+\hat{m}_k(\boldsymbol{\lambda})\boldsymbol{\beta}_k\|$$

To bound $\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + \hat{m}_k(\boldsymbol{\lambda})\boldsymbol{\beta}_k\|$, we use the definition of $\hat{m}_k(\boldsymbol{\lambda})$:

$$\begin{aligned}
& \left\| y - \sum_{j=1}^{J} g_j(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + \hat{m}_j(\boldsymbol{\lambda})\boldsymbol{\beta}_j) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j \left(P_j \left(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + \hat{m}_k(\boldsymbol{\lambda})\boldsymbol{\beta}_k \right) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + \hat{m}_k(\boldsymbol{\lambda})\boldsymbol{\beta}_k\|^2 \right) \\
\leq & \frac{1}{2} \|y - \sum_{j=1}^{J} g(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j})\|_T^2 + \sum_{j=1}^{J} \lambda_j \left(P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}\|_2^2 \right) \\
= & \frac{1}{2} \|y - \sum_{j=1}^{J} g(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j})\|_T^2 + \sum_{j=1}^{J} \lambda_j^{(1)} \left(P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}\|_2^2 \right) + \sum_{j=1}^{J} (\lambda_j - \lambda_j^{(1)}) \left(P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}\|_2^2 \right) \\
\leq & C + J\lambda_{max} \max_{j=1:J} \left(P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}\|_2^2 \right)
\end{aligned}$$

To bound the term $\max_{j=1:J} \left(P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}\|_2^2 \right)$, we use the basic inequality for $\hat{\boldsymbol{\theta}}_{\lambda^{(1)}}$:

$$\begin{aligned}
\sum_{j=1}^{J} \lambda_j^{(1)} \left(P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}\|_2^2 \right) & \leq \frac{1}{2} \|y - \sum_{j=1}^{J} g(\cdot|\hat{\boldsymbol{\theta}}_j^*)\|_T^2 + \sum_{j=1}^{J} \lambda_j^{(1)} \left(P_j(\hat{\boldsymbol{\theta}}_j^*) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_j^*\|_2^2 \right) \\
& \leq C
\end{aligned}$$

Since

$$\lambda_{min} \left(\max_{j=1:J} P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}\|_2^2 \right) \leq \sum_{j=1}^{J} \lambda_j^{(1)} \left(P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}\|_2^2 \right)$$

then we have that

$$\max_{j=1:J} P_j(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j}\|_2^2 \leq \frac{C}{\lambda_{min}}$$

Therefore

$$\frac{1}{2} \left\| y - \sum_{j=1}^{J} g_j(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + \hat{m}_j(\boldsymbol{\lambda})\boldsymbol{\beta}_j) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j \left(P_j \left(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + \hat{m}_k(\boldsymbol{\lambda})\boldsymbol{\beta}_k \right) + \frac{w}{2} \|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + \hat{m}_k(\boldsymbol{\lambda})\boldsymbol{\beta}_k\|^2 \right) \leq \left(1 + \frac{J\lambda_{max}}{\lambda_{min}} \right) C$$

This implies that

$$\|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + m_\ell \boldsymbol{\beta}_k\| \leq \sqrt{\left(1 + \frac{J\lambda_{max}}{\lambda_{min}} \right) \frac{2C}{\lambda_{min} w}} \tag{8}$$

and

$$\left\| y - \sum_{j=1}^{J} g_j(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + \hat{m}_j(\boldsymbol{\lambda})\boldsymbol{\beta}_j) \right\|_T \leq \sqrt{\left(1 + \frac{J\lambda_{max}}{\lambda_{min}}\right) 2C} \tag{9}$$

If combine the assumption (5) with (8), we get

$$\left| \frac{\partial}{\partial m_k} P_k(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + m_k\boldsymbol{\beta}_k) + w\langle \boldsymbol{\beta}_k, \hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + m_k\boldsymbol{\beta}_k \rangle \right|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})} \leq \|\boldsymbol{\beta}_k\| \left(K_k + w\sqrt{\left(1 + \frac{J\lambda_{max}}{\lambda_{min}}\right) \frac{2C}{\lambda_{min} w}} \right)$$

On the other hand, suppose the other assumption (6) is satisfied. Then we will need to use the implicit differentiation equation (7). Rearranging, we get

$$\begin{aligned} \frac{\partial}{\partial m_k} \left(P_k(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + m_k\boldsymbol{\beta}_k) \right) \Big|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})} + w\langle \boldsymbol{\beta}_k, \hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + m_k\boldsymbol{\beta}_k \rangle &= \frac{1}{\lambda_k} \left\langle \frac{\partial}{\partial m} g_k \left(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + m_k\boldsymbol{\beta}_k \right) \Big|_{m=\hat{m}(\lambda)}, y - \sum_{j=1}^{J} g_j(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + m_j\boldsymbol{\beta}_j) \right\rangle_T \\ &\leq \frac{1}{\lambda_{min}} K_k \|\boldsymbol{\beta}_k\| \left\| y - \sum_{j=1}^{J} g_j \left(\cdot|\hat{\boldsymbol{\theta}}_{\lambda^{(1)},j} + \hat{m}_j(\boldsymbol{\lambda})\boldsymbol{\beta}_j \right) \right\|_T \end{aligned}$$

Plugging in (9), we get

$$\left| \frac{\partial}{\partial m_k} P_k(\hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + m_k\boldsymbol{\beta}_k) + w\langle \boldsymbol{\beta}_k, \hat{\boldsymbol{\theta}}_{\lambda^{(1)},k} + m_k\boldsymbol{\beta}_k \rangle \right|_{\boldsymbol{m}=\hat{\boldsymbol{m}}(\boldsymbol{\lambda})} \leq \|\boldsymbol{\beta}_k\| \frac{1}{\lambda_{min}} K_k \sqrt{\left(1 + \frac{J\lambda_{max}}{\lambda_{min}}\right) 2C}$$

Using these upper bounds, we can bound D_3 by the diagonal matrix

$$\left\{ \max_{k=1,...,J} d_k \right\} diag\left(\{\|\beta_k\|\}_{k=1}^{J} \right) \succeq D_3$$

where for $k = 1, ..., J$

$$d_k = \begin{cases} \left(K_k + w\sqrt{\left(1 + \frac{J\lambda_{max}}{\lambda_{min}}\right) \frac{2C}{\lambda_{min} w}} \right) & \text{if assumption (5) holds for } k \\ \frac{1}{\lambda_{min}} K_k \sqrt{\left(1 + \frac{J\lambda_{max}}{\lambda_{min}}\right) 2C} & \text{if assumption (6) holds for } k \end{cases}$$

3. We bound the norm of $\nabla_\lambda \hat{m}_k(\lambda)$ for all $k = 1, ..., J$.

For every $k = 1, ..., J$, we have

$$\begin{aligned}
\|\nabla_\lambda \hat{m}_k(\lambda)\| &= \|Me_k\| \\
&= \|D_3 \left(S + D_1 + D_2\right)^{-1} e_k\| \\
&\leq \left\{\max_{k=1,...,J} d_k\right\} \left\|diag\left(\{\|\beta\|_k\}_{k=1}^J\right)\left(S + D_1 + D_2\right)^{-1} e_k\right\| \\
&\leq \left\{\max_{k=1,...,J} d_k\right\} \max_\ell \|\boldsymbol{\beta}_\ell\| \left\|\left(S + D_1 + D_2\right)^{-1} e_k\right\| \\
&\leq \left\{\max_{k=1,...,J} d_k\right\} \max_\ell \|\boldsymbol{\beta}_\ell\| \left\|D_2^{-1} e_k\right\| \qquad (10)
\end{aligned}$$

The last line follows from the matrix inverse lemma: Since $S + D_1$ is a PSD matrix, then

$$\left\|\left(S + D_1 + D_2\right)^{-1} e_k\right\| \leq \left\|D_2^{-1} e_k\right\|$$

Now consider (8) for

$$k \coloneqq \ell_{max} = \arg\max_\ell \|\boldsymbol{\beta}_\ell\|$$

(Notice we can choose any k in $1, ..., J$. The inequality holds for all k so we just choose the k that is most interesting for our problem.)
We have

$$\begin{aligned}
\|\nabla_\lambda \hat{m}_{\ell_{max}}(\lambda)\| &\leq \left\{\max_{k=1,...,J} d_k\right\} \|\boldsymbol{\beta}_{\ell_{max}}\| \left\|D_2^{-1} e_{\ell_{max}}\right\| \\
&= \left\{\max_{k=1,...,J} d_k\right\} \|\boldsymbol{\beta}_{\ell_{max}}\| \lambda_{\ell_{max}}^{-1} w^{-1} \|\boldsymbol{\beta}_{\ell_{max}}\|_2^{-2} \\
&\leq \left\{\max_{k=1,...,J} d_k\right\} \|\boldsymbol{\beta}_{\ell_{max}}\|^{-1} \lambda_{min}^{-1} w^{-1}
\end{aligned}$$

4. Apply the Mean Value Theorem
Since the training criterion is smooth, then $\hat{m}_{\ell_{max}}(\lambda)$ is a continuous, differentiable function.
By the MVT, we have that there exists an $\alpha \in (0, 1)$ such that

$$\begin{aligned}
\left|\hat{m}_{\ell_{max}}(\boldsymbol{\lambda}^{(2)}) - \hat{m}_{\ell_{max}}(\boldsymbol{\lambda}^{(1)})\right| &= \left|\left\langle \boldsymbol{\lambda}^{(2)} - \boldsymbol{\lambda}^{(1)}, \nabla_\lambda \hat{m}_{\ell_{max}}(\boldsymbol{\lambda})\right\rangle_{\boldsymbol{\lambda} = \alpha\boldsymbol{\lambda}^{(1)} + (1-\alpha)\boldsymbol{\lambda}^{(2)}}\right| \\
&\leq \left\|\boldsymbol{\lambda}^{(2)} - \boldsymbol{\lambda}^{(1)}\right\| \left\{\max_{k=1,...,J} d_k\right\} \lambda_{min}^{-1} w^{-1} \|\boldsymbol{\beta}_{\ell_{max}}\|^{-1}
\end{aligned}$$

We know that $\hat{m}_k(\boldsymbol{\lambda}^{(2)}) - \hat{m}_k(\boldsymbol{\lambda}^{(1)}) = \mathbf{1}$ for all $k = 1, .., J$. Rearranging the inequality above, we get

$$\max_k \|\boldsymbol{\theta}_{\lambda^{(1)},k} - \boldsymbol{\theta}_{\lambda^{(2)},k}\| = \|\boldsymbol{\beta}_{\ell_{max}}\| \leq \left\|\boldsymbol{\lambda}^{(2)} - \boldsymbol{\lambda}^{(1)}\right\| \left\{\max_{k=1,...,J} d_k\right\} \lambda_{min}^{-1} w^{-1}$$

3 Nonsmooth Penalties

Suppose we are dealing with parametric regression problems from Section 1 or 2. We keep all the same assumptions, except those that concern the smoothness of the penalties.

Recall that $\Lambda \subseteq \mathbb{R}^J$. Consider the measure space over Λ with respect to the Lebesgue measure μ. We suppose that for a given dataset (X, y), suppose the following three assumptions hold:

Assumption (1): Let the penalized training criterion be denoted $L_T(\boldsymbol{\theta}, \boldsymbol{\lambda})$. Denote the differentiable space of $L_T(\cdot, \boldsymbol{\lambda})$ at any point $\boldsymbol{\theta}$ as

$$\Omega^{L_T(\cdot,\boldsymbol{\lambda})}(\boldsymbol{\theta}) = \left\{\boldsymbol{\eta} | \lim_{\epsilon \to 0} \frac{L_T(\boldsymbol{\theta} + \epsilon\boldsymbol{\eta}) - L_T(\boldsymbol{\theta})}{\epsilon} \text{ exists}\right\}$$

Suppose there is a set $\Lambda_{smooth} \subseteq \Lambda$ such that $\mu\left(\Lambda_{smooth}^C\right) = 0$ and for every $\boldsymbol{\lambda} \in \Lambda_{smooth}$, there exists a ball with nonzero radius centered at $\boldsymbol{\lambda}$, denoted $B(\boldsymbol{\lambda})$, such that the following conditions hold:

Cond 1: For all $\boldsymbol{\lambda}' \in B(\boldsymbol{\lambda})$, the training criterion $L_T(\cdot, \cdot)$ is twice differentiable along directions in $\Omega^{L_T(\cdot,\cdot)}\left(\hat{\boldsymbol{\theta}}_\lambda\right)$. (So technically the twice-differentiable space is constant)

Cond 2: $\Omega^{L_T(\cdot,\boldsymbol{\lambda})}\left(\hat{\boldsymbol{\theta}}_\lambda\right)$ is a local optimality space of $B(\boldsymbol{\lambda})$:

$$\arg\min_{\theta \in \Theta} L_T\left(\boldsymbol{\theta}, \boldsymbol{\lambda}'\right) = \arg \min_{\theta \in \Omega^{L_T(\cdot,\boldsymbol{\lambda})}\left(\hat{\boldsymbol{\theta}}_\lambda\right)} L_T\left(\boldsymbol{\theta}, \boldsymbol{\lambda}'\right) \ \forall \boldsymbol{\lambda}' \in B(\boldsymbol{\lambda})$$

Cond 3: (Not necessary if we keep the ridge penalty) There is an orthonormal basis U_λ of $\Omega^{L_T(\cdot,\boldsymbol{\lambda})}\left(\hat{\boldsymbol{\theta}}_\lambda\right)$ such that the Hessian of the training criterion taken along directions U_λ is invertible.

Assumption (2): For every $\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}} \in \Lambda_{smooth}$, let the line segment between the two points be denoted

$$\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}) = \left\{\alpha\boldsymbol{\lambda^{(1)}} + (1-\alpha)\boldsymbol{\lambda^{(2)}} : \alpha \in [0,1]\right\}$$

Suppose the intersection $\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}) \cap \Lambda_{smooth}^C$ is countable.

Assumption (3): All the conditions specified in Section 1 and 2 that bound the spectrum of P_j or g_j only need to apply when the directional derivatives exist. That is, the condition on the spectrum of the penalty derivative is now

$$\left|\frac{\partial}{\partial m} P_j\left(\boldsymbol{\theta} + m\boldsymbol{\beta}\right)\right| \leq K\|\boldsymbol{\beta}\|_2 \text{ if } \frac{\partial}{\partial m} P_j\left(\boldsymbol{\theta} + m\boldsymbol{\beta}\right) \text{exists}$$

Similarly, we would change the condition on the spectrum of the function derivative to

$$\left|\frac{\partial}{\partial m} g_j\left(\boldsymbol{\theta}+m\boldsymbol{\beta}\right)\right| \leq K\|\boldsymbol{\beta}\|_2 \text{ if } \frac{\partial}{\partial m} g_j\left(\boldsymbol{\theta}+m\boldsymbol{\beta}\right) \text{exists}$$

Under these assumptions, the same Lipschitz conditions hold for dataset (X, y) and every $\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}} \in \Lambda_{smooth}$.

Proof

Consider any $\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}} \in \Lambda_{smooth}$. The length of $\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$ covered by set A can be expressed as

$$\mu_1\left(A \cap \mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})\right)$$

where μ_1 is the Lebesgue measure over the line segment $\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$. (So if $A \cap \mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$ is just a line segment, it is the length $\|A \cap \mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})\|_2$)

By the Differentiability Cover Lemma below, there exists a countable set of points $\cup_{i=1}^{\infty} \boldsymbol{\ell}^{(i)} \subset \mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$ such that the union of their "balls of differentiabilities" entirely cover $\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$:

$$\max_{\{\boldsymbol{\ell}^{(i)}\}_{i=1}^{\infty}} \mu_1\left(\cup_{i=1}^{\infty} B(\boldsymbol{\ell}^{(i)}) \cap \mathcal{L}\left(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}\right)\right) = \left\|\mathcal{L}\left(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}\right)\right\|_2$$

Let

$$\left\{\boldsymbol{\ell}_{max}^{(i)}\right\}_{i=1}^{\infty} = \left\{\arg\max_{\{\boldsymbol{\ell}^{(i)}\}} \mu_1\left(\cup_{i=1}^{\infty} B(\boldsymbol{\ell}^{(i)}) \cap \mathcal{L}\left(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}\right)\right)\right\} \cup \left\{\boldsymbol{\lambda}^{(1)}, \boldsymbol{\lambda}^{(2)}\right\}$$

Let P be the intersections of the boundary of $B\left(\boldsymbol{\ell}_{max}^{(i)}\right)$ with the line segment $\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$:

$$P = \cup_{i=1}^{\infty} \text{Bd}B\left(\boldsymbol{\ell}_{max}^{(i)}\right) \cap \mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$$

Every point $p \in P$ can be expressed as $\alpha_p \boldsymbol{\lambda^{(1)}} + (1-\alpha_p)\boldsymbol{\lambda^{(2)}}$ for some $\alpha_p \in [0, 1]$. This means we can order these points $\{\boldsymbol{p}^{(i)}\}_{i=1}^{\infty}$ by increasing α_p. By our assumptions, the differentiable space of the training criterion must be constant over the interior of line segment $\mathcal{L}\left(\boldsymbol{p^{(i)}}, \boldsymbol{p^{(i+1)}}\right)$ (so there might be bad behavior at the endpoints). Let the differentiable space over the interior of line segment $\mathcal{L}\left(\boldsymbol{p^{(i)}}, \boldsymbol{p^{(i+1)}}\right)$ be denoted Ω_i.

By our assumptions, the differentiable space is also a local optimality space. Let $U^{(i)}$ be an orthonormal basis of Ω_i. For each i, we can express $\hat{\boldsymbol{\theta}}_\lambda$ for all $\boldsymbol{\lambda} \in \text{Int}\left\{\mathcal{L}\left(\boldsymbol{p^{(i)}}, \boldsymbol{p^{(i+1)}}\right)\right\}$ as

$$\hat{\boldsymbol{\theta}}_\lambda = U^{(i)}\hat{\boldsymbol{\beta}}_\lambda$$

$$\hat{\boldsymbol{\beta}}_\lambda = \arg\min_{\beta} L_T(U^{(i)}\boldsymbol{\beta}, \boldsymbol{\lambda})$$

Now apply the result in Section 1 or 2 over every line segment $\mathcal{L}\left(\boldsymbol{p^{(i)}}, \boldsymbol{p^{(i+1)}}\right)$. To do this, we must modify the proofs to take directional derivatives along the columns of $U^{(i)}$. We can establish that there is a constant $c > 0$ independent of i such that for all $i = 1, 2...$, we have

$$\left\|\hat{\boldsymbol{\beta}}_{p^{(i)}} - \hat{\boldsymbol{\beta}}_{p^{(i+1)}}\right\|_2 \leq c\|\boldsymbol{p^{(i)}} - \boldsymbol{p^{(i+1)}}\|_2$$

Finally, we can sum these inequalities. By the triangle inequality,

$$\begin{aligned}
\left\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(1)}}} - \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda^{(2)}}}\right\|_2 &\leq \sum_{i=1}^{\infty} \|\hat{\boldsymbol{\theta}}_{p^{(i)}} - \hat{\boldsymbol{\theta}}_{p^{(i+1)}}\|_2 \\
&= \sum_{i=1}^{\infty} \|U^{(i)}\hat{\boldsymbol{\beta}}_{p^{(i)}} - U^{(i)}\hat{\boldsymbol{\beta}}_{p^{(i+1)}}\|_2 \\
&= \sum_{i=1}^{\infty} \|\hat{\boldsymbol{\beta}}_{p^{(i)}} - \hat{\boldsymbol{\beta}}_{p^{(i+1)}}\|_2 \\
&\leq \sum_{i=1}^{\infty} c\|\boldsymbol{p^{(i)}} - \boldsymbol{p^{(i+1)}}\|_2 \\
&= c\|\boldsymbol{\lambda}^{(1)} - \boldsymbol{\lambda}^{(2)}\|_2
\end{aligned}$$

Lemma - Differentiability Cover

For any $\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}} \in \Lambda_{smooth}$, there exists a countable set of points $\cup_{i=1}^{\infty}\boldsymbol{\ell}^{(i)} \subset \mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$ such that the union of their "balls of differentiabilities" entirely cover $\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$

$$\max_{\{\boldsymbol{\ell}^{(i)}\}_{i=1}^{\infty}} d_{\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}}\left(\cup_{i=1}^{\infty} B(\boldsymbol{\ell}^{(i)})\right) = \left\|\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})\right\|$$

Proof

We prove this by contradiction. Let

$$\left\{\boldsymbol{\ell}_{max}^{(i)}\right\}_{i=1}^{\infty} = \arg\max_{\{\boldsymbol{\ell}^{(i)}\}_{i=1}^{\infty}} d_{\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}}\left(\cup_{i=1}^{\infty} B(\boldsymbol{\ell}^{(i)})\right)$$

and for contradiction, suppose that the covered length is less than the length of the line segment:

$$d_{\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}}\left(\cup_{i=1}^{\infty} B(\boldsymbol{\ell}_{max}^{(i)})\right) < \left\|\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})\right\|$$

By assumption (2), since $\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}) \cap \Lambda^C_{smooth}$ is countable, there must exist a point $p \in \mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}) \setminus \left\{\cup_{i=1}^{\infty} B(\boldsymbol{\ell}^{(i)}_{max})\right\}$ such that $p \notin \Lambda^C_{smooth}$. However if we consider the set of points $\left\{\boldsymbol{\ell}^{(i)}_{max}\right\}_{i=1}^{\infty} \cup \{p\}$, then

$$d_{\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}} \left(\cup_{i=1}^{\infty} B(\boldsymbol{\ell}^{(i)}_{max})\right) < d_{\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}}} \left(\cup_{i=1}^{\infty} B(\boldsymbol{\ell}^{(i)}_{max}) \cup B(p)\right)$$

This is a contradiction of the definition of $\{\boldsymbol{\ell}^{(i)}_{max}\}$. Therefore we should always be able to cover $\mathcal{L}(\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}})$ with "balls of differentiability."

4 Example

4.1 Penalties that satisfy the conditions

We will show penalties that satisfy the condition

$$\frac{\partial}{\partial m} P(\boldsymbol{\theta} + m\boldsymbol{\beta}) \leq K \|\boldsymbol{\beta}\|_2$$

for some constant $K > 0$.

Ridge:

The perturbation isn't necessary if there is already a ridge penalty in the original penalized regression problem. Just set the penalties $P_j(\boldsymbol{\theta}) \equiv 0$ and fix $w = 2$.

Lasso:

$$\begin{aligned} \frac{\partial}{\partial m} \|\theta + m\beta\|_1 &= \langle sgn(\theta + m\beta), \beta \rangle \\ &\leq \|sgn(\theta + m\beta)\|_2 \|\beta\|_2 \\ &\leq p\|\beta\|_2 \end{aligned}$$

so $K = p$ in this case.

Generalized Lasso: let G be the maximum eigenvalue of D.

$$\begin{aligned} \frac{\partial}{\partial m} \|D(\theta + m\beta)\|_1 &= \langle sgn(D(\theta + m\beta)), D\beta \rangle \\ &\leq \|sgn(D(\theta + m\beta))\|_2 \|D\beta\|_2 \\ &\leq pG\|\beta\|_2 \end{aligned}$$

so $K = pG$ in this case.

Group Lasso:

If we have un-pooled penalty parameters as follows

$$\sum_{j=1}^{J} \lambda_j \|\boldsymbol{\theta}^{(j)} + m^{(j)}\boldsymbol{\beta}^{(j)}\|_2$$

then we need the following bound for every $j = 1, ..., J$

$$\begin{aligned} \frac{\partial}{\partial m^{(j)}} \|\boldsymbol{\theta}^{(j)} + m^{(j)}\boldsymbol{\beta}^{(j)}\|_2 &= \left\langle \frac{\boldsymbol{\theta}^{(j)} + m^{(j)}\boldsymbol{\beta}^{(j)}}{\|\boldsymbol{\theta}^{(j)} + m^{(j)}\boldsymbol{\beta}^{(j)}\|_2}, \boldsymbol{\beta}^{(j)} \right\rangle \\ &\leq \|\boldsymbol{\beta}^{(j)}\|_2 \end{aligned}$$

So $K = 1$ in this case.
If there is a single penalty parameter for the entire group laso penalty as follows

$$\lambda \sum_{j=1}^{J} \|\boldsymbol{\theta}^{(j)} + m\boldsymbol{\beta}^{(j)}\|_2$$

then

$$\begin{aligned} \frac{\partial}{\partial m} \sum_{j=1}^{J} \|\boldsymbol{\theta}^{(j)} + m\boldsymbol{\beta}^{(j)}\|_2 &= \sum_{j=1}^{J} \left\langle \frac{\boldsymbol{\theta}^{(j)} + m\boldsymbol{\beta}^{(j)}}{\|\boldsymbol{\theta}^{(j)} + m\boldsymbol{\beta}^{(j)}\|_2}, \boldsymbol{\beta}^{(j)} \right\rangle \\ &\leq \sum_{j=1}^{J} \|\boldsymbol{\beta}^{(j)}\|_2 \\ &\leq \sqrt{J}\|\boldsymbol{\beta}\|_2 \end{aligned}$$

and $K = \sqrt{J}$.

4.2 Sobolev

Given a function h, the Sobolev penalty for h is

$$P(h) = \int (h^{(r)}(x))^2 dx$$

The Sobolev penalty is used in nonparametric regression models, but such nonparametric regression models can be re-expressed in parametric form. We will use this to understand the smoothness of models fitted in this manner.

Consider the class of smoothing splines

$$\left\{ \hat{g}(\cdot|\lambda) = \arg\min_{g \in \mathcal{G}} \frac{1}{2} \left\| y - \sum_{j=1}^{J} g_j(x_j) \right\|_T^2 + \sum_{j=1}^{J} \lambda_j P(g_j) : \lambda \in \Lambda \right\}$$

Each function $\hat{g}_j(\cdot|\lambda)$ is a spline that can be expressed as the weighted sum of B normalized B-splines of degree $r+1$ for a given set of knots:

$$\hat{g}_j(x|\lambda) = \sum_{i=1}^{B} \theta_i N_{j,i}(x)$$

Note that the normalized B-splines have the property that they sum up to one at all points within the boundary of the knots. Also recall that B-splines are non-negative.

Therefore we can re-express the class of smoothing splines as a set of function parameters

$$\left\{ \hat{\boldsymbol{\theta}}_\lambda = \arg\min_{\theta} \frac{1}{2} \left\| y - \sum_{j=1}^{J} N_{T,j} \boldsymbol{\theta}_j \right\|_T^2 + \sum_{j=1}^{J} \lambda_j P_j(\boldsymbol{\theta}_j) : \lambda \in \Lambda \right\}$$

where $N_{T,j}$ is a matrix of the evaluations of the normalized B-spline basis at x_j. $P_j(\boldsymbol{\theta_j})$ is the Sobolev penalty and can be written as $\boldsymbol{\theta}_j^T V_j \boldsymbol{\theta}_j$ for an appropriate penalty matrix V_j. We will not need to express anything in terms of V_j so the penalty will be just written as $P_j(\boldsymbol{\theta}_j)$.

Instead of considering the original smoothing spline problem with the roughness penalty, we will add a ridge penalty on the function parameters

$$\left\{ \hat{\boldsymbol{\theta}}_\lambda = \arg\min_{\theta} \frac{1}{2} \left\| y - \sum_{j=1}^{J} N_{T,j} \boldsymbol{\theta}_j \right\|_T^2 + \sum_{j=1}^{J} \lambda_j \left(P_j(\boldsymbol{\theta}_j) + \frac{w}{2} \|\boldsymbol{\theta}_j\|_2^2 \right) : \lambda \in \Lambda \right\}$$

Let

$$C = \frac{1}{2} \left\| y - \sum_{j=1}^{J} N_{T,j} \boldsymbol{\theta}_j^* \right\|_T^2 + \lambda_{max} \sum_{j=1}^{J} \left(P_j(\boldsymbol{\theta}_j^*) + \frac{w}{2} \|\boldsymbol{\theta}_j^*\|_2^2 \right)$$

Then for any $\boldsymbol{\lambda^{(1)}}, \boldsymbol{\lambda^{(2)}} \in \Lambda$ we have for all $j = 1, ..., J$

$$\|\boldsymbol{\theta}_{\lambda^{(1)},j} - \boldsymbol{\theta}_{\lambda^{(2)},j}\|_2 \leq \left\| \boldsymbol{\lambda}^{(2)} - \boldsymbol{\lambda}^{(1)} \right\|_2 \lambda_{min}^{-1} w^{-1} \left(\frac{1}{\lambda_{min}} B \sqrt{\left(1 + \frac{J\lambda_{max}}{\lambda_{min}} \right) \frac{2C}{\lambda_{min} w}} \right)$$

Moreover,

$$\left\| \sum_{j=1}^{J} \hat{g}_j(x_j|\boldsymbol{\lambda}^{(1)}) - \hat{g}_j(x_j|\boldsymbol{\lambda}^{(2)}) \right\|_\infty \leq \left\| \boldsymbol{\lambda}^{(2)} - \boldsymbol{\lambda}^{(1)} \right\|_2 J\sqrt{B}\lambda_{min}^{-1} w^{-1} \left(\frac{1}{\lambda_{min}} B \sqrt{\left(1 + \frac{J\lambda_{max}}{\lambda_{min}}\right) \frac{2C}{\lambda_{min} w}} \right)$$

Proof

To apply the result from Section 2, we just need to bound the spectral norm

$$\|\nabla_\theta g_j(X_{T,j}|\boldsymbol{\theta})\| = \|N_{T,j}\|$$

Note that the eigenvalue of $N_{T,j}$ is bounded by B since the maximum eigenvalue of a non-negative matrix is bounded by its maximum row sum. In the case of $N_{T,j}$, since it is the values of normalized B-splines, each row is at most the number of B-spline basis functions. That is, we have for all $j = 1, ..., J$

$$\|\nabla_\theta g_j(X_{T,j}|\boldsymbol{\theta})\| = \|N_{T,j}\| \leq B$$

Hence for all $\boldsymbol{\theta}, \boldsymbol{\beta}, m'$, we have

$$\left\| \frac{\partial}{\partial m} g_j(X_{T,j}|\boldsymbol{\theta} + m\boldsymbol{\beta}) \Big|_{m=m'} \right\| \leq B\|\boldsymbol{\beta}\|$$

Apply the result from Section 2 to get the result

$$\left\| \boldsymbol{\theta}_{\lambda^{(1)},j} - \boldsymbol{\theta}_{\lambda^{(2)},j} \right\|_2 \leq \left\| \boldsymbol{\lambda}^{(2)} - \boldsymbol{\lambda}^{(1)} \right\|_2 \lambda_{min}^{-1} w^{-1} \left(\frac{1}{\lambda_{min}} B \sqrt{\left(1 + \frac{J\lambda_{max}}{\lambda_{min}}\right) \frac{2C}{\lambda_{min} w}} \right)$$

The "moreover" statement follows from the fact that for any point $\boldsymbol{x}$, we have

$$\begin{aligned}
\left|\sum_{j=1}^{J} \hat{g}_j(x_j|\boldsymbol{\lambda}^{(1)}) - \hat{g}_j(x_j|\boldsymbol{\lambda}^{(2)})\right| &= \left|\sum_{j=1}^{J}\sum_{i=1}^{B}\left(\hat{\theta}_{\boldsymbol{\lambda}^{(1)},j,i} - \hat{\theta}_{\boldsymbol{\lambda}^{(2)},j,i}\right) N_{j,i}(x_j)\right| \\
&\leq \sum_{j=1}^{J}\sum_{i=1}^{B}\left|\left(\hat{\theta}_{\boldsymbol{\lambda}^{(1)},j,i} - \hat{\theta}_{\boldsymbol{\lambda}^{(2)},j,i}\right) N_{j,i}(x_j)\right| \\
&\leq \sum_{j=1}^{J}\sum_{i=1}^{B}\left|\hat{\theta}_{\boldsymbol{\lambda}^{(1)},j,i} - \hat{\theta}_{\boldsymbol{\lambda}^{(2)},j,i}\right| \\
&\leq \sum_{j=1}^{J}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)},j} - \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(2)},j}\|_1 \\
&\leq \sqrt{B}\sum_{j=1}^{J}\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)},j} - \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(2)},j}\|_2
\end{aligned}$$

where the second inequality uses the fact that normalized B-splines have value at most 1. Therefore

$$\left\|\sum_{j=1}^{J} \hat{g}_j(x_j|\lambda^{(1)}) - \hat{g}_j(x_j|\lambda^{(2)})\right\|_\infty \leq \sqrt{B}\sum_{j=1}^{J}\left\|\hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(1)},j} - \hat{\boldsymbol{\theta}}_{\boldsymbol{\lambda}^{(2)},j}\right\|_2$$